

# Agent Workflow

## Table of contents

<b>1</b>	<b>Agent Workflow</b>	<b>1</b>
1.1	The Agent's Core Loop . . . . .	1
1.2	System Prompt and Agent Persona . . . . .	2
1.3	Persistence and Memory . . . . .	2
1.4	Example Interaction Flow . . . . .	3

## 1 Agent Workflow

This document explains the internal workflow of the `code-agent` and how it processes your requests. Understanding this flow can help you interact with the agent more effectively and troubleshoot issues.

### 1.1 The Agent's Core Loop

The `code-agent` operates as an interactive loop, driven by a LangGraph `StateGraph`. This graph orchestrates the interaction between the user, the Language Model (LLM), and the agent's tools.

Here's a simplified overview of the workflow for each user query:

1. **User Input:** You provide a natural language query to the agent.
2. **Context Retrieval:** The agent first performs a similarity search against its persistent conversation history ( stored in `ChromaDB`). Relevant past interactions are retrieved and added to the current context.
3. **LLM Invocation (Reasoning):** The LLM receives the user's query along with the retrieved context and a set of tool descriptions. Its task is to decide the next best action:

- **Use a Tool:** If the LLM determines that a tool is needed to fulfill the request (e.g., creating a file, reading code, generating tests), it outputs a `ToolCall` specifying the tool’s name and arguments.
  - **Direct Response:** If the LLM can answer the query directly without needing a tool, it generates a natural language response.
4. **Tool Execution:** If a `ToolCall` is made, the specified tool is executed. The tool performs its action (e.g., creates a file, modifies code, runs a script) and returns its output.
  5. **LLM Re-invocation (Tool Output Processing):** The LLM receives the output from the executed tool. It then uses this new information to decide the next step:
    - **Further Tool Use:** If the goal is not yet achieved, the LLM might call another tool or the same tool with different arguments.
    - **Final Response:** Once the LLM determines the user’s request has been fulfilled, it generates a final natural language response.
  6. **History Update:** Both the user’s query and the agent’s final response (including any tool outputs) are stored in the persistent `ChromaDB` vector store, enriching the agent’s memory for future interactions.

This loop continues until the agent provides a final response to the user.

## 1.2 System Prompt and Agent Persona

The agent’s behavior and decision-making are heavily influenced by its **system prompt**. This prompt defines the agent’s persona, its capabilities, and its instructions for interacting with the user and using its tools.

The `code-agent` is instructed to act as a proactive and autonomous “senior software engineer.” It is encouraged to use its tools directly to achieve user goals without constantly asking for explicit instructions or claiming it cannot perform a task if a tool is available.

## 1.3 Persistence and Memory

The `code-agent` maintains conversational memory across sessions using a local `ChromaDB` vector store.

- **Location:** The memory is stored in a `.code_agent_memory` directory within your project’s root.
- **Mechanism:** Each user query and agent response is embedded and added to this vector store. When a new query arrives, the agent retrieves semantically similar past interactions to provide relevant context to the LLM.
- **Benefits:** This allows the agent to remember past tasks, learn from previous interactions, and maintain a more coherent conversation flow over time.

## 1.4 Example Interaction Flow

Consider the user query: “Create a new Python file named `hello.py` and add a function `greet()` that prints ‘Hello, World!’”

1. **User Input:** “Create a new Python file named `hello.py` and add a function `greet()` that prints ‘Hello, World!’”
2. **LLM Reasoning:** The LLM analyzes the request and determines that the `new-file` tool is appropriate. It formulates the `file_path` and `content` arguments.
3. **Tool Execution:** The `new-file` tool is called with `file_path='hello.py'` and `content='def greet():\n print("Hello, World!")'`.
4. **Tool Output:** The tool successfully creates the file and returns a message like “Successfully created `hello.py`”.
5. **LLM Final Response:** The LLM receives the tool’s output, confirms the task is complete, and generates a final response to the user: “I have created `hello.py` with the `greet()` function.”
6. **History Update:** The entire interaction (user query, tool call, tool output, agent response) is saved to ChromaDB.

This iterative process allows the agent to break down complex requests into manageable steps, execute them, and provide a comprehensive response.